

Introduction AI/ML

Main Areas of Artificial Intelligence



from wikipedia

- **computer vision**

data: spatial structures (e.g., images), SOTA: Convolutional Neural Networks (CNN)

- **natural language processing**

data: sequential structures (e.g., text), SOTA: transformers

- **automated decision making, robotics**

data: sequential actions (e.g., games), SOTA: reinforcement learning

agency:

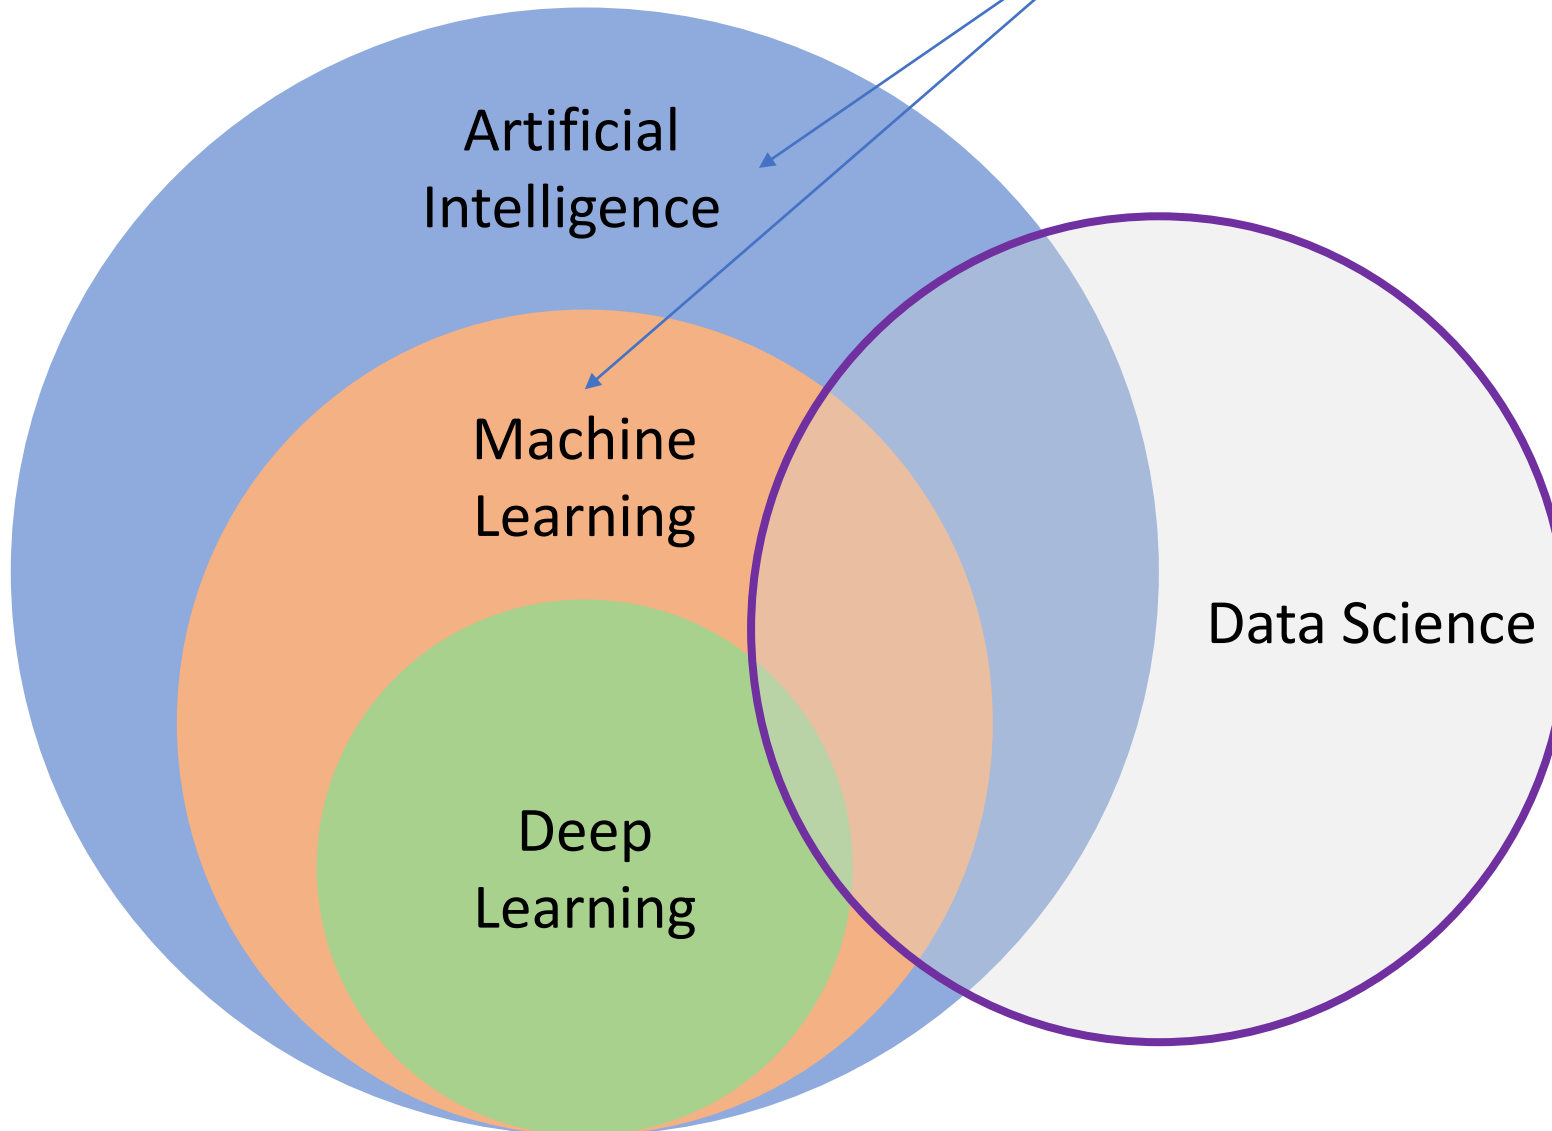
perception – thought – action

All of these are enabled by one key ingredient:

learning from experience/data (**Machine Learning**)

data can also be tabular (structured): columns as features, rows as independent samples

blend of diverse components from different domains
(statistics, optimization, computer science, ...)



Deep Learning:

special kind of ML
algorithms using *deep*
neural networks (e.g.,
CNNs, transformers)

Data Science:

extract knowledge from
data (by means of ML,
among other things)

ML: Learning from Experience/Data

mainly exploiting statistical dependencies with the aim of **generalization** to new (e.g., future) data (compare with human reasoning by [analogies](#))

training (usually offline optimization):

ML algorithm + data = explicit algorithm (to be used at inference time)

→ reduction of complexity and much better generalizability compared to handcrafted algorithms

analogy: Humans do not hit the ground running (storage capacity of DNA limited) but have learning capabilities.

Ladder of Generalization

shallow learning:

representation encoded in features


→ feature engineering

deep learning:

representation encoded in network

→ feature/representation learning


(hierarchy of concepts learned from raw data in deep graph with many layers)



When to Use ML (= Learning from Data)

automation

too complex for rules




from wikipedia

object recognition, chat bot, ...

complexity / uncertainty

too complex for humans



AlphaFold

protein structure predictions, demand forecasting, ...

more scientific use cases: medicine (imaging, diagnosis, drug design), particle physics (analysis of collider experiments), material science (material properties and design of new materials), ...

MACHINE LEARNING

training target available
(labeled or past data)

SUPERVISED

CLASSIFICATION



REGRESSION

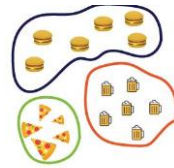


learning by teacher
(high-dimensional curve fitting)

data not labeled
in any way

UNSUPERVISED

CLUSTERING

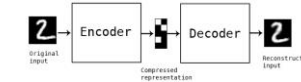
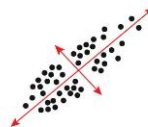


GENERATIVE MODELS

ASSOCIATION



DIMENSIONALITY REDUCTION



learning by observation
(pattern recognition)

no supervision, but goal-based
interaction with environment

REINFORCEMENT LEARNING

LEARN STATE OR ACTION VALUES

LEARN POLICY DIRECTLY



learning by trial-and-error
(sequential decision making)

unsupervised and reinforcement learning can
both be cast as supervised-learning setup

Supervised Learning

learning by teacher → usually rather narrow tasks (passive approach)

Target Quantity

- **known in training:** labeled samples or observations from past
- to be **predicted** for unknown cases (e.g., future values)

Features

input information that is

- correlated to target quantity
- known at prediction time




Example: Spam Filtering

Classify emails as spam or no spam

use accordingly **labeled**
emails as training set

use information like
occurrence of specific
words or email length
as **features**

features x_1 and x_2
spam, no spam



But Before: Data Processing

environment:

[WSL](#), [Python](#), [virtualenv](#), [pip](#)

scientific Python stack:

[NumPy](#), [pandas](#), [matplotlib](#)

coding example: stock market data